

# Machine Learning

In this jupyter notebook, we are training a machine learning model on the data that we exported from Matlab, with our features. We split the data into the 2 recordings, so that leakage was reduced - the first recording was used for the training dataset, and the second recording was used for testing. We compared usage of a Random Forest Classifier and Support Vector Machine. Our result is that the Random Forest Classifier has a higher accuracy then the support vector machine - therefore we will use it over the Support Vector Machine.

## Importing the Data

```
In [79]: import pandas as pd
```

```
In [80]: df = pd.read_csv("ml_data.csv")
df.head()
```

```
Out[80]: avg_bpm    rmssd    sdnn    gsr_avg    mean_absolute_derivative    std_gs
          0    71.287129  0.046690  0.046231  9339.129032                8.688751  552.34609
          1    88.757396  0.012000  0.008485  8033.696062                10.214207 438.87837
          2    84.905660  0.061400  0.053454  8002.003634                7.397931  410.89007
          3    84.905660  0.061400  0.053454  7980.096348                7.805495  386.54236
          4    81.190798  0.063676  0.078017  7962.944878                5.007646  365.67888
```

```
In [81]: import sklearn.model_selection as ms
X = df.loc[:, df.columns != "stress"]
X.head()
```

```
Out[81]: avg_bpm    rmssd    sdnn    gsr_avg    mean_absolute_derivative    std_gs
          0    71.287129  0.046690  0.046231  9339.129032                8.688751  552.34609
          1    88.757396  0.012000  0.008485  8033.696062                10.214207 438.87837
          2    84.905660  0.061400  0.053454  8002.003634                7.397931  410.89007
          3    84.905660  0.061400  0.053454  7980.096348                7.805495  386.54236
          4    81.190798  0.063676  0.078017  7962.944878                5.007646  365.67888
```

```
In [82]: Y = df.loc[:, df.columns == "stress"]
Y.head()
```

Out[82]: **stress**

<b>0</b>	0
<b>1</b>	0
<b>2</b>	0
<b>3</b>	0
<b>4</b>	0

```
In [83]: from sklearn.model_selection import train_test_split

# split it so that the first recording is training dataset, and 2nd record
# data comes like this, all calm recordings 1 and 2, then all stress recordings
# We want all calm recordings 1 and all stress recordings 1 in X_train
from collections import Counter
import pandas as pd
from sklearn.utils import shuffle

print(Counter(Y["stress"]))

calm_recording_1_length = int(1168/4)
calm_recording_2_length = int(1200/4)
stress_recording_1_length = int(1284/4)
stress_recording_2_length = int(1520/4)

# Define slice indices
calm1_start = 0
calm1_end = calm1_start + calm_recording_1_length

calm2_start = calm1_end
calm2_end = calm2_start + calm_recording_2_length

stress1_start = calm2_end
stress1_end = stress1_start + stress_recording_1_length

stress2_start = stress1_end
stress2_end = stress2_start + stress_recording_2_length

X_train = pd.concat([
    X.iloc[calm1_start:calm1_end, :],
    X.iloc[stress1_start:stress1_end, :]
])

X_test = pd.concat([
    X.iloc[calm2_start:calm2_end, :],
    X.iloc[stress2_start:stress2_end, :]
])

Y_train = pd.concat([
    Y.iloc[calm1_start:calm1_end],
    Y.iloc[stress1_start:stress1_end]
])

Y_test = pd.concat([
    Y.iloc[calm2_start:calm2_end],
    Y.iloc[stress2_start:stress2_end]
```

```
] )

X_train, Y_train = shuffle(X_train, Y_train, random_state=42)
X_test, Y_test = shuffle(X_test, Y_test, random_state=42)

print(X_train.shape, X_test.shape)
print(Y_train.shape, Y_test.shape)
```

```
Counter({1: 699, 0: 590})
```

```
(613, 6) (676, 6)
```

```
(613, 1) (676, 1)
```

```
In [84]: X_train.head()
```

```
Out[84]:    avg_bpm      rmssd      sdnn      gsr_avg  mean_absolute_derivative      std_
670    77.888360   0.098752   0.119164  3359.891429           1.531244   348.858
101    69.870845   0.088546   0.113002  4828.571973           2.764526   824.900
131    66.683132   0.146319   0.149301  6124.757865           2.462128   396.242
693    85.376532   0.048758   0.053750  2257.462345           1.308668   381.997
909    72.727273   0.106623   0.103216   754.154143           2.022058   433.567
```

```
In [85]: Y_train.head()
```

```
Out[85]:    stress
670      1
101      0
131      0
693      1
909      1
```

## Training the SVM (Support Vector Machine)

```
In [86]: from sklearn.svm import SVC
model = SVC()
model
```

Out[86]:

▼ SVC ⓘ ⓘ

## ► Parameters

C	1.0
kernel	'rbf'
degree	3
gamma	'scale'
coef0	0.0
shrinking	True
probability	False
tol	0.001
cache_size	200
class_weight	None
verbose	False
max_iter	-1
decision_function_shape	'ovr'
break_ties	False
random_state	None

In [87]:

```
import numpy as np
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
model.fit(X_train_scaled, np.ravel(Y_train))
```

Out[87]:

▼ SVC ⓘ ⓘ

► Parameters

C	1.0
kernel	'rbf'
degree	3
gamma	'scale'
coef0	0.0
shrinking	True
probability	False
tol	0.001
cache_size	200
class_weight	None
verbose	False
max_iter	-1
decision_function_shape	'ovr'
break_ties	False
random_state	None

In [88]: model.score(X\_test\_scaled, Y\_test)

Out[88]: 0.8683431952662722

In [89]: from sklearn.metrics import classification\_report, confusion\_matrix

```
predictions = model.predict(X_test_scaled)
print(classification_report(Y_test, predictions, zero_division=0))
```

	precision	recall	f1-score	support
0	0.77	1.00	0.87	298
1	1.00	0.76	0.87	378
accuracy			0.87	676
macro avg	0.89	0.88	0.87	676
weighted avg	0.90	0.87	0.87	676

In [90]: from sklearn.model\_selection import GridSearchCV  
from sklearn.pipeline import Pipeline

```
pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('svm', SVC())
])
```

```
param_grid = {'svm__C': [0.1, 1, 10, 100, 1000],
              'svm__gamma': [1, 0.1, 0.01, 0.001, 0.0001, 'scale'],
              'svm__kernel': ['rbf']}
grid = GridSearchCV(pipeline, param_grid, verbose=0)
```

```
In [91]: grid.fit(X_train, np.ravel(Y_train))
```

```
Out[91]: Pipeline(steps=[('scaler', StandardScaler()),
                         ('svm', SVC())])
```

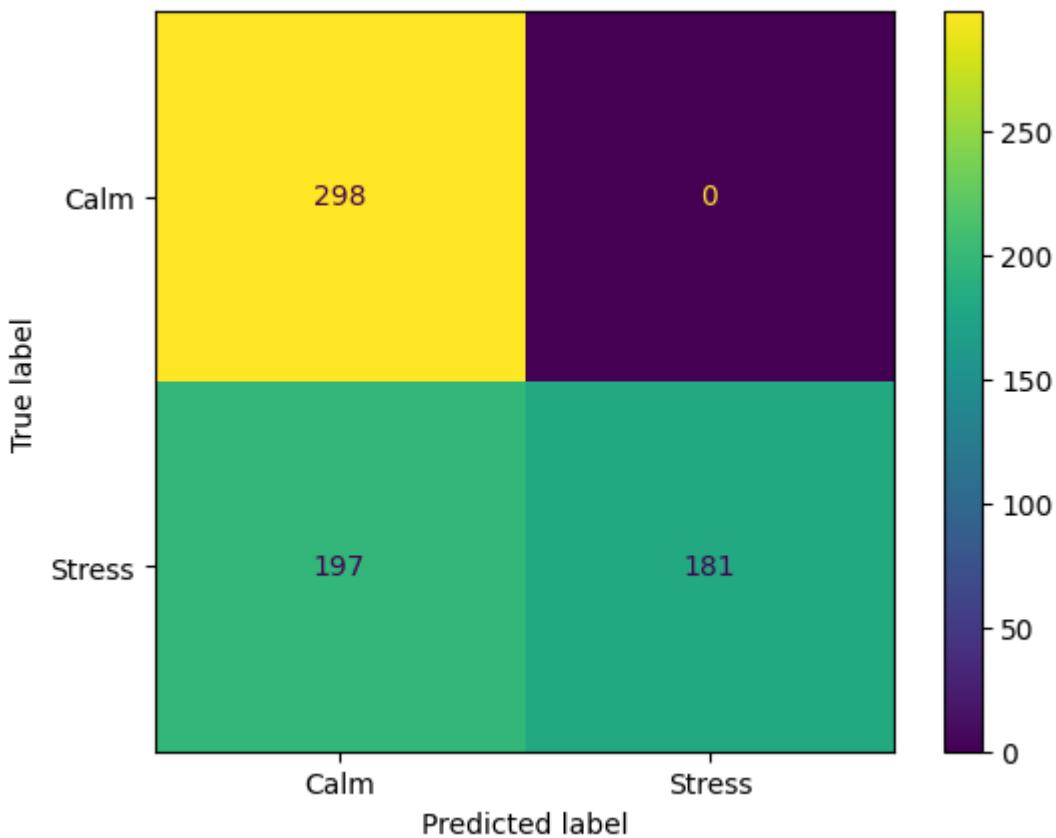
```
In [92]: print(grid.best_params_)
print(grid.best_estimator_)
print(grid.best_score_)
```

```
{'svm__C': 0.1, 'svm__gamma': 1, 'svm__kernel': 'rbf'}
Pipeline(steps=[('scaler', StandardScaler()), ('svm', SVC(C=0.1, gamma=1))])
1.0
```

```
In [93]: grid_predictions = grid.predict(X_test)
print(grid.score(X_test, Y_test))
print(classification_report(Y_test, grid_predictions, zero_division=0))
```

```
0.7085798816568047
precision    recall   f1-score   support
          0       0.60      1.00      0.75      298
          1       1.00      0.48      0.65      378
   accuracy                           0.71      676
    macro avg       0.80      0.74      0.70      676
 weighted avg       0.82      0.71      0.69      676
```

```
In [94]: from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt
cm = confusion_matrix(Y_test, grid_predictions)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Calm",
                                                               "Agitated"])
disp.plot()
plt.show()
```



## Results of SVM

Our results for SVM was an accuracy 71%

## Training Random Forest Classifier

```
In [95]: from sklearn.ensemble import RandomForestClassifier  
rf = RandomForestClassifier()
```

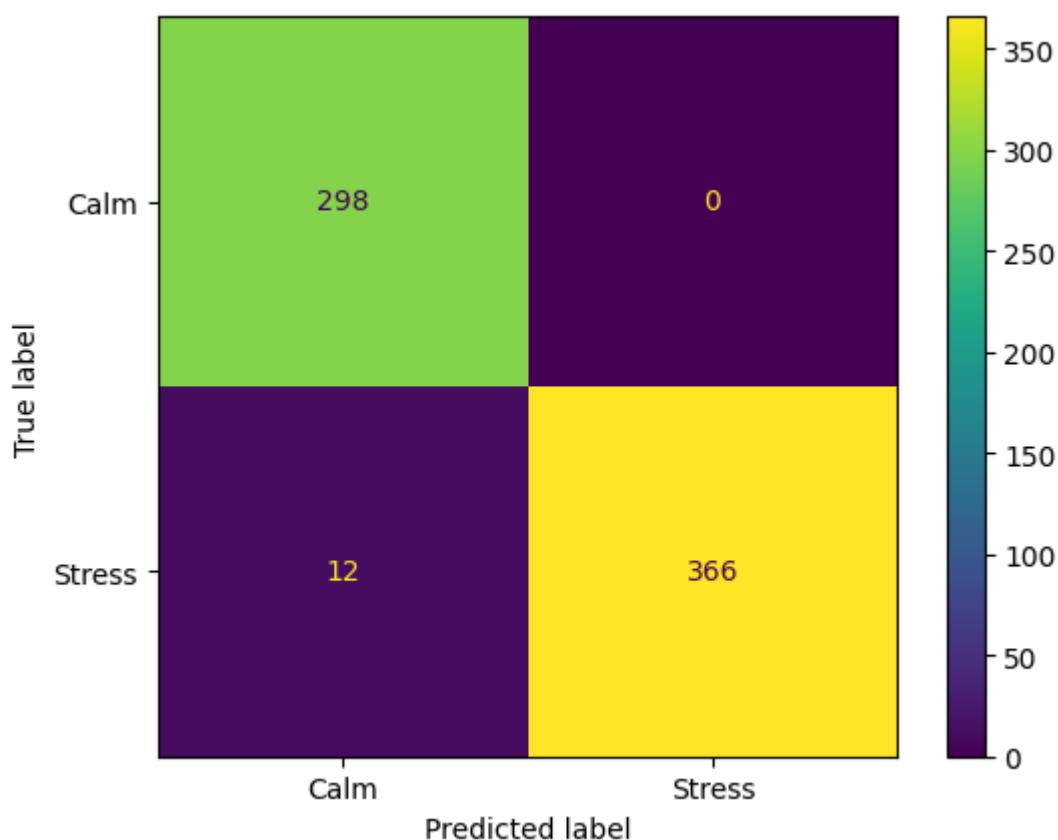
```
In [96]: param_grid_rf = {  
    'max_depth': [3,5,7,10],  
    'n_estimators': [100, 200, 300, 400, 500],  
    'max_features': [10, 20, 30 , 40],  
    'min_samples_leaf': [1, 2, 4]  
}  
  
grid_rf = GridSearchCV(rf, param_grid_rf, scoring='f1', n_jobs=-1, verbose=1)  
grid_rf.fit(X_train, np.ravel(Y_train))
```

```
Out[96]:  
    ▶ GridSearchCV ⓘ ⓘ  
    ▶ best_estimator_: RandomForestClassifier  
        ▶ RandomForestClassifier ?
```

```
In [97]: grid_rf.score(X_test, Y_test)
```

```
Out[97]: 0.9838709677419355
```

```
In [98]: cm_rf = confusion_matrix(Y_test, grid_rf.predict(X_test))
disp = ConfusionMatrixDisplay(cm_rf, display_labels=["Calm", "Stress"])
disp.plot()
plt.show()
```



## Results

Random Forest Classifier has a far better score at 98% compared to Support Vector Machine's 71%. Therefore we will use a Random Forest Classifier over the Support Vector Machine for this dataset, due to the improved performance.

```
In [99]: import joblib
joblib.dump(grid_rf, "model.pkl")
```

```
Out[99]: ['model.pkl']
```